

Stanford University, CS 106B, Autumn 2013
Homework Assignment 2: Word Ladder and N-Grams
due Wednesday, October 16, 2013, 2:00pm

Thanks to Julie Zelenski and Jerry Cain for the assignment idea and much of the text; Random Writer comes from Joe Zachary.

The purpose of this assignment is to practice using collections. You will use vector, stack, queue, set, and map. Since it is hard to craft a single problem that exercises all of these collections, this is a two-part assignment.

Files:

We will provide you with two separate **ZIP archives**, each of which contains a starter version of that part of your project. You should download these archives from the class web site and write the rest of the code. For **Part A** (Word Ladders), you will turn in only the following file:

- **wordladder.cpp**, the C++ code for the Word Ladder program (including a **main** function)

For **Part B** (N-Grams), you will turn in the following files:

- **ngrams.cpp**, the C++ code for the N-grams program (including a **main** function)
- **myinput.txt**, your own unique input file representing text to read in as your program's input

The ZIP archives contain other files and libraries; you should not modify them. When grading your code, we will run your file with our own original versions of the support files, so your code must work with them.

Part A: Word Ladders

A **word ladder** is a connection from one word to another formed by changing one letter at a time with the constraint that at each step the sequence of letters still forms a valid word. For example, here is a word ladder connecting the word "code" to the word "data". Each changed letter is underlined as an illustration:

code → cade → cate → date → data

There are many other word ladders that connect these two words, but ours is the shortest. That is, there might be others of the same length, but none with fewer steps than this one.

In the first part of this assignment, you will write a program that finds a minimum-length word ladder between two words typed by the user. Your code must use the **Stack** and **Queue** collections from Chapter 5, along with a particular provided algorithm to find the shortest such sequence. This part is shorter and simpler than Part B.

Here is an example **log of interaction** between your program and the user (with console input underlined):

```
Welcome to CS 106B Word Ladder.
If you give me two English words, I will transform the
first into the second by changing one letter at a time.

Please type two words: code data
Ladder from data back to code:
data date cate cade code
Have a nice day.
```

Notice that the word ladder prints out in reverse order, from the second word back to the first. If there are multiple valid word ladders of the same length between a given starting and ending word, your program would not need to generate exactly the ladder shown in this log, but you must generate one of minimum length.

You may assume valid input. For example, you may assume that the user types exactly two words and that both of them are valid words in the English dictionary, and that they are not the same word. You may also assume that the English dictionary file exists and is readable by your program. If invalid input occurs, your program's behavior is unspecified; it can do whatever you want, including crashing.

You will need to keep a **dictionary** of all English words. We provide a file named **dictionary.txt** that contains these words, one per line. Read this file as input and choose an efficient collection to store and look up words. Note that you should not ever need to loop over the dictionary as part of solving this problem.

Part A Implementation Details:

Finding a word ladder is a specific instance of a **shortest-path problem** of finding a path from a start position to a goal. Shortest-path problems come up in routing Internet packets, comparing gene mutations, and so on. The strategy we will use for finding a shortest path is called **breadth-first search ("BFS")**, a search process that expands out from a start position, considering all possibilities that are one step away, then two steps away, and so on, until a solution is found. BFS guarantees that the first solution you find will be as short as any other. (*Breadth-first search is not the most efficient algorithm for generating minimal word ladders, but later we will touch on improved search algorithms, and in advanced courses such as CS161 you will learn even more efficient alternatives.*)

For word ladders, start by examining ladders that are one step away from the original word, where only one letter is changed. Then check all ladders that are two steps away, where two letters have been changed. Then three, four, etc. We implement the breadth-first algorithm using a **queue** to store partial ladders that represent possibilities to explore. Each partial ladder is a **stack**, which means that your overall collection is a **queue of stacks**.

Here is a partial **pseudocode** description of the algorithm to solve the word-ladder problem:

```
function wordLadder(w1, w2):
    create an empty queue of stacks.
    create/add a stack containing {w1} to the queue.
    while the queue is not empty:
        dequeue the partial-ladder stack from the front of the queue.
        if the word on top of the stack is the destination word:
            hooray! output the elements of the stack as the solution.
        otherwise:
            for each valid English word that is a "neighbor" (differs by 1 letter)
              of the word on top of the stack:
                if that neighbor word has not already been used in a ladder before:
                    create a copy of the current ladder stack.
                    put the neighbor word on top of the copy stack.
                    add the copy stack to the end of the queue.
```

Some of the pseudocode corresponds almost one-to-one with actual C++ code. One part that is more abstract is the part that instructs you to examine each **"neighbor"** of a given word. A neighbor of a given word *w* is a word of the same length as *w* that differs by exactly 1 letter from *w*. For example, **"date"** and **"date"** are neighbors.

It is *not* appropriate to look for neighbors by looping over the entire dictionary every time; this is way too slow. To find all neighbors of a given word, use two nested loops: one that goes through each character index in the word, and one that loops through the letters of the alphabet from **a-z**, replacing the character in that index position with each of the 26 letters in turn. For example, when examining neighbors of **"date"**, you'd try:

- **aate, bate, cate, ..., zate** ← all possible neighbors where only the 1st letter is changed
- **date, dbte, dcte, ..., dzte** ← all possible neighbors where only the 2nd letter is changed
- ...
- **data, datb, datc, ..., datz** ← all possible neighbors where only the 4th letter is changed

Note that many of the possible words along the way (**aate, dbte, datz**, etc.) are not valid English words. Your algorithm has access to an **English dictionary**, and each time you generate a word using this looping process, you should look it up in the dictionary to make sure that it is actually a legal English word.

Another subtle issue is that you **do not reuse words** that have been included in a previous ladder. For example, suppose that you have add the partial ladder **cat** → **cot** → **cog** to the queue. Later on, if your code is processing ladder **cat** → **cot** → **con**, one neighbor of **con** is **cog**, so you might want to examine **cat** → **cot** → **con** → **cog**.

But doing so is unnecessary. If there is a word ladder that begins with these four words, then there must be a shorter one that, in effect, cuts out the middleman by eliminating the unnecessary word **con**. As soon as you've enqueued a ladder ending with a specific word, you've found a minimum-length path from the starting word to the end word in the ladder, so you *never* have to enqueue that end word again.

To implement this strategy, keep track of the set of words that have already been used in *any* ladder. Ignore those words if they come up again. Keeping track of what words you've used also eliminates the possibility of getting trapped in an infinite loop by building a circular ladder, such as **cat** → **cot** → **cog** → **bog** → **bag** → **bat** → **cat**.

Part B: N-Grams

In the second part of this assignment, you will write a program that reads an input file and uses it to build a large data structure of word groups called "N-grams" as a basis for randomly generating new text that sounds like it came from the same author as that file. You will use the **Map** and **Vector** collections from Chapter 5.

At right is an example **log of interaction** between your program and the user (console input underlined).

But what, you may ask, is an N-gram?

```
Welcome to CS 106B Random Writer (aka 'N-Grams').
This program makes random text based on a document.
Give me an input file and an 'N' value for groups
of words, and I'll generate random text for you.
```

```
Input file? hamlet.txt
Value of N? 3
```

```
# of random words to generate (0 to quit)? 40
... chapel. Ham. Do not believe his tenders, as you
go to this fellow. Whose grave's this, sirrah?
Clown. Mine, sir. [Sings] O, a pit of clay for to
the King that's dead. Mar. Thou art a scholar; speak
to it. ...
```

```
# of random words to generate (0 to quit)? 20
... a foul disease, To keep itself from noyance; but
much more handsome than fine. One speech in't I
chiefly lov'd. ...
```

```
# of random words to generate (0 to quit)? 0
Exiting.
```

The "Infinite Monkey Theorem" states that an infinite number of monkeys typing random keys forever would eventually produce the works of William Shakespeare. That's silly, but could a monkey randomly produce a *new work* that "sounded like" Shakespeare's works, with similar vocabulary, wording, punctuation, etc.? What if we chose *words* at random, instead of individual letters? Suppose that rather than each word having an equal probability of being chosen, we weighted the probability based on how often that word appeared in Shakespeare's works?

Picking random words would likely produce gibberish, but let's look at *chains of two words* in a row. For example, perhaps Shakespeare uses the word "to" occurs 10 times total, and 7 of those occurrences are followed by "be", 1 time by "go", and 2 times by "eat". We can use those ratios when choosing the next word. If the last word we chose is "to", randomly choose "be" next 7/10 of the time, "go" 1/10 of the time, and "eat" 2/10. We never choose any other word to follow "to". We call a chain of two words like this, such as "to be", a *2-gram*.

```
+-----+
| Chose "to". |----> choose "be"   (7/10 chance)
| Next random word? |----> choose "go"   (1/10 chance)
+-----+----> choose "eat"   (2/10 chance)
```

Go, get you have seen, and now he makes as itself? (2-gram)

A sentence of 2-grams isn't great, but look at chains of 3 words (*3-grams*). If we chose the words "to be", what word should follow? If we had a collection of all sequences of 3 words-in-a-row with probabilities, we could make a weighted random choice. If Shakespeare uses "to be" 22 times and follows them with "or" 5 times, "in" 3 times, "with" 10 times, and "alone" 4 times, we could use these weights to randomly choose the next word. So now the algorithm would pick the third word based on the first two, and the fourth based on the (second+third), and so on.

```
+-----+----> choose "or"   ( 5/22 chance)
| Chose {"to", "be"}. |----> choose "in"   ( 3/22 chance)
| Next random word? |----> choose "with" (10/22 chance)
+-----+----> choose "alone" ( 4/22 chance)
```

One woe doth tread upon another's heel, so fast they follow. (3-gram)

You can generalize the idea from 2-grams to N-grams for any integer N. If you make a collection of all groups of N words along with their possible following words, you can use this to select an N+1'th word given the preceding N words. The higher N level you use, the more similar the new random text will be to the original data source. Here is a random sentence generated from 5-grams of *Hamlet*, which is starting to sound a lot like the original:

I cannot live to hear the news from England, But I do prophesy th' election lights on Fortinbras. (5-gram)

Each particular piece of text randomly generated in this way is also called a *Markov chain*. Markov chains are very useful in computer science and elsewhere, such as artificial intelligence, machine learning, economics, and statistics.

Part B, Algorithm Step 1: Building Map of N-Grams

In this program, you will read the input file one word at a time and build a particular compound collection, a **map** from prefixes to suffixes. If you are building 3-grams, that is, N -grams for $N=3$, then your code should examine sequences of 2 words and look at what third word follows those two. For later lookup, your map should be built so that it connects a collection of $N-1$ words with another collection of all possible suffixes; that is, all possible N 'th words that follow the previous $N-1$ words in the original text. For example if you are computing N -grams for $N=3$ and the pair of words "to be" is followed by "or" twice and "just" once, your collection should map the key {to, be} to the value {or, just, or}. The following figure illustrates the map you should build from the file:

When reading the input file, the idea is to keep a window of $N-1$ words at all times, and as you read each word from the file, discard the first word from your window and append the new word. The following figure shows the file being read and the map being built over time as each of the first few words is read to make 3-grams:

to be [^] or not to be just ...	map = {} window = {to, be}
to be or [^] not to be just ...	map = {{to, be} : {or}} window = {be, or}
to be or not [^] to be just ...	map = {{to, be} : {or}, {be, or} : {not}} window = {or, not}
to be or not to [^] be just ...	map = {{to, be} : {or}, {be, or} : {not}, {or, not} : {to}} window = {not, to}
to be or not to be [^] just ...	map = {{to, be} : {or}, {be, or} : {not}, {or, not} : {to}, {not, to} : {be}} window = {to, be}
to be or not to be just [^] ...	map = {{to, be} : {or, just}, {be, or} : {not}, {or, not} : {to}, {not, to} : {be}} window = {be, just}
...	...
to be or not to be just be who you want to be or not okay you want okay	map = {{to, be} : {or, just, or}, {be, or} : {not, not}, {or, not} : {to, okay}, {not, to} : {be}, {be, just} : {be}, {just, be} : {who}, {be, who} : {you}, {who, you} : {want}, {you, want} : {to, okay}, {want, to} : {be}, {not, okay} : {you}, {okay, you} : {want}, {want, okay} : {to}, {okay, to} : {be}}
input file, tiny.txt	resulting map of 3-gram suffixes

Note that the **order matters**: For example, the prefix {you, are} is different from the prefix {are, you}. Note that the same word can occur multiple times as a suffix, such as "or" occurring twice after the prefix {to, be}.

Also notice that the map **wraps around**. For example, if you are computing 2-grams, perform 2 more iterations to connect the last 2 prefixes in the end of the file to the first 2 words at the start of the file. In our example above, this leads to {want, okay} connecting to "to" and {okay, to} connecting to "be". If we were doing 5-grams, we would perform 5 more iterations and connect to the first 5 words in the file, and so on. This turns out to be very useful to help your algorithm later on in the program.

Your **should not change case or strip punctuation** of words as you read them. The casing and punctuation turns out to help the sentences start and end in a more authentic way. Just store the words in the map as you read them.

Part B, Algorithm Step 2: Generating Random Text

To generate random text from your map of N -grams, first choose a random starting point for the document. To do this, pick a randomly chosen key from your map. Each key is a collection of $N-1$ words. Those $N-1$ words will form the start of your random text. This collection of $N-1$ words will be your **sliding "window"** as you create your text.

For all subsequent words, use your map to look up all possible next words that can follow the current $N-1$ words, and randomly choose one with appropriate weighted probability. If you have built your map the way we described, as a map from $\{\text{prefix}\} \rightarrow \{\text{suffixes}\}$, this simply amounts to choosing one of the possible suffix words at random. Once you have chose your random suffix word, slide your current "window" of $N-1$ words by discarding the first word in the window and appending the new suffix. The following diagram illustrates the text generation algorithm.

Action(s)	Current ($N-1$) "window"	Output so far
choose a random start	{"who", "you"}	who you
choose new word; shift	{"you", "want"}	who you want
choose new word; shift	{"want", "okay"}	who you want okay
choose new word; shift	{"okay", "to"}	who you want okay to
...

Note that in our random example, at one point our window was `{want, okay}`. This was the end of the original input file. Nothing actually follows that prefix, which is why it was important that we made our map **wrap around** from the end of the file to the start, so that if our window ever ends up at the last $N-1$ words from the document, we won't get stuck unable to generate further random text.

Since your random text likely won't happen to start and end at the beginning/end of a sentence, just prefix and suffix your random text with "..." to indicate this. Here is another partial log of execution:

```
Input file? tiny.txt
Value of N? 3
# of random words to generate (0 to quit)? 16
... who you want okay to be who you want to be or not to be or ...
```

Assume valid input. For example, you may assume that the user types the name of a file that exists and is readable, that the user will type a value for N that is a positive integer no greater than the number of words in the file, that the input file will be at least N words in length, and that the number of random words to generate will be at least $(N - 1)$.

Development Strategy and Hints:

This program can be tricky if you don't develop and debug it step-by-step. Don't try to write everything all at once. Make sure to **test** each part of the algorithm before you move on. See the **Homework FAQ** for more tips.

- Think about exactly what **types of collections** to use for each part. Are duplicates allowed? Does order matter? Do you need random access? Where will you add/remove elements? Etc. Note that some parts of each program require you to make compound collections, that is, a collection of collections.
- Test each function with a very **small input** first. For example, use input file `tiny.txt` with a small number of words because you can **print your entire map** and examine its contents.
- Recall that you can **print** the contents of any collection to `cout` and examine its contents for debugging.
- Remember that when you assign one collection to another using the `=` operator, it makes a full copy of the entire contents of the collection. This could be useful if you want to copy a collection.
- To choose a random prefix from a map, consider using the map's `keys` member function, which returns a **Vector** containing all of the keys in the map. For **randomness** in general, include `"random.h"` and call the global function `randomInteger(min, max)`.
- You can loop over the elements of a vector or set using a **foreach** loop. A **foreach** also works on a map, iterating over the keys in the map. You can look up each associated value based on the key in the loop.
- Don't forget to test your input on **unusual inputs**, like large and small values of N , large/small # of words to generate, large and small input files, and so on. It's hard to verify random input, but you can look in smallish input files to verify that a given word really does follow a given prefix from your map.

Implementation and Grading:

All items mentioned in the "Implementation and Grading" from the previous assignment(s) specs also apply here. Please refer to those documents as needed. Note the instructions in the previous assignment about procedural decomposition, variables, types, parameters, value vs. reference, and commenting. Don't forget to **cite any sources** you used in your comments. Refer to the course **Style Guide** for a more thorough discussion of good coding style.

Collections: Additionally, on this assignment part of your Style grade comes from making intelligent decisions about what kind of **collections** from the Stanford C++ library to use at each step of your algorithm, as well as using those collections elegantly. As much as possible, **pass collections by reference**, because passing them by value makes an expensive copy of the collection. Do not use pointers, arrays, or STL containers on this program.

Don't forget to use the course web site's **Output Comparison Tool** to help check your output for various test cases.

Honor Code: Please remember to follow the **Honor Code** when working on this assignment. Submit your own work and do not look at others' solutions. Also please do not give out your solution and do not place a solution to this assignment on a public web site or forum. If you need help, please seek out our available resources to help you.

For reference, our solution to Part A (the part you need to write) is around 40 lines long including spacing and comments. Our solution to Part B (the complete file) is around 130 lines long, and it has 4 functions besides **main**, though you don't need to match these numbers to get full credit; they are just here as a ballpark or sanity check.

Copyright © Stanford University and Marty Stepp, licensed under Creative Commons Attribution 2.5 License. All rights reserved.